

Relating Electric Vehicle Ownership to Household Power Consumption

Scott Sullivan

September 25, 2016

1 Introduction

What happens when someone plugs in their electric car? The amount of power drawn from an electric vehicle (EV) can be higher than the peak usage of a typical household. Therefore, when someone plugs in their EV, the usage should noticeably spike. By looking for spikes in household usage, it should be possible to infer properties of the data, such as which houses own EVs and when those EVs are charging.

By using techniques from statistics and machine learning, it was possible to predict which houses own EVs with an accuracy of $(84 \pm 2)\%$, and when those houses were charging with an accuracy of $(93.0 \pm .6)\%$. This paper describes at a high-level the techniques used to make these predictions. Section II describes the Markov Model used to predict when an EV was charging. Section III describes the K-Nearest-Neighbors (KNN) approach used to predict which houses own EVs. Section IV describes other interesting features found in the data. For more details, see the accompanying IPython Notebooks.

2 When an EV is Charging

In order to determine when an EV was charging, a Markov Model was constructed with hidden variables given by $P(\text{charging})$ and $P(\text{not charging})$, and observation variables given by the derivatives in the usage at each timestep. The derivatives were found by using a cubic spline at each point in the data, and the first four derivatives were used as 4-Dimensional input (including the usage itself, the zeroth derivative). Before the Markov Model was trained, each household's usage was normalized by subtracting its mean and dividing by its standard deviation. Then the four sets of derivatives were normalized with each other by subtracting their overall mean and dividing by their deviation, in order to equalize the weights of each dimension.

The transition probabilities were calculated by separating each pair of adjacent timesteps into four groups: not charging \rightarrow not charging (00), not charging \rightarrow charging (01), charging \rightarrow not charging (10), and charging \rightarrow charging (11). Then the derivatives at the second step in each transition were used to build a pair of 4D classifiers ($0 \rightarrow n$, $1 \rightarrow n$), which were used to predict $P(\text{charge}_i | \text{charge}_{i-1}, D_i)$, where i was the timestep, and D were the derivatives. A plot showing the results of the Markov Model is shown in

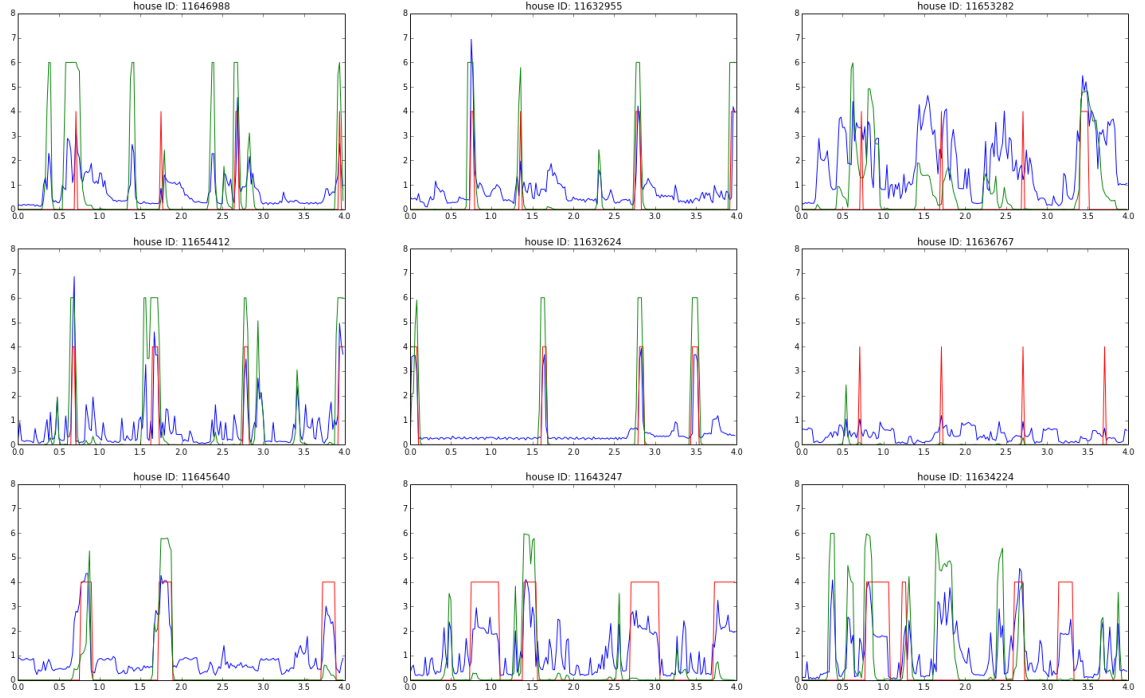


Figure 1: Blue curve: Household usage (kWh/30-min interval), as a function of time (day). Red curve: Times at which an EV is charging, scaled to 4 for visual clarity. Green curve: Probability an EV is charging as predicted by the Markov Model, scaled to 6 for visual clarity.

figure 1. The overall accuracy of this method was $(93.0 \pm .06)\%$. For more details, see Part II - Predicting When an EV is Charging.ipynb.

3 Which Houses Own EVs

In order to determine which houses own EVs, a K-Nearest-Neighbor (KNN) classifier was constructed by calculating an array of statistical features for each household and experimentally testing which feature (or pair of features) contained the most predictive power. The features included:

- The average usage
- The standard deviation in the usage

- An average of the absolute values of the zeroth, first, second, and third derivatives
- Same as above, but increasing the influence of larger values by first taking the square, cube, 4th power, and 5th power of each point in the timeseries before averaging.
- Same as above, but after normalizing each usage graph first (as in section 2).

In addition, each feature was normalized with itself so they would all have the same weight.

A 1D KNN classifier was built for each feature, as well as a 2D KNN classifier for all possible pairs of features. Because trying such a large number of different models was likely to result in overfitting, the data was split into train, validation, and test data. The validation data was used to select the most accurate 1D KNN and the most accurate 2D KNN. The test data was then used to gauge the true accuracy of each model by running once for the most accurate 1D KNN and once for the most accurate 2D KNN.

The most accurate classifier was the 1D KNN for average first derivative, weighted with $k=5$, with no household normalization. The most accurate 1D KNN was actually more accurate than the most accurate 2D KNN. It's likely all the pairs of features were not truly independent, causing models based on pairs of features to lose accuracy if the second statistic was not as powerful as the first. A histogram of this feature is shown in figure 2. The overall accuracy of this method was $(84 \pm 2)\%$. For more details see Part III - Predicting Which Houses Own EVs.ipynb.

4 Other Interesting Features in the Data

During the analysis, several noteworthy features were found in the dataset which are described below. The first was how much power an EV consumed on average. To determine EV power consumption, first the base usage of each house was calculated by averaging all the usages when the EV was not charging. This was then subtracted from each usage during charging, and the resulting usages were averaged for each house. (The ends of each charge were also cut off, to remove times which were not charging for the complete interval). It turns out the EV power consumption had two peaks, as seen

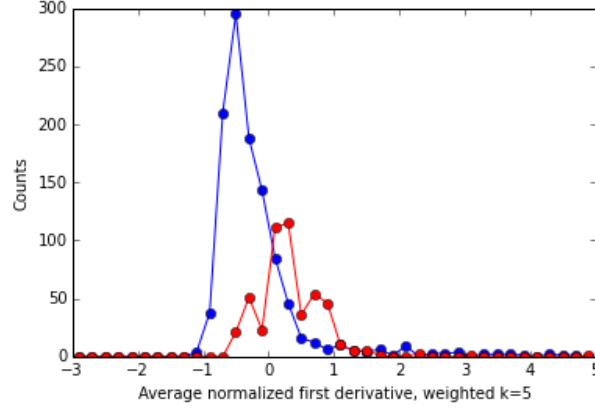


Figure 2: Histogram of $\text{average}(|\text{usage first derivative}|^5)$, separated into households without EVs (blue) and with EVs (red).

in figure 3. It seems that there are two models of EV, one which consumes about 3.5 kW when plugged in, and one which consumes about 6.5 kW.

The other interesting feature was the usage for each interval averaged across households. A day/night cycle is clearly visible, as shown in figure 4. For more details, see Part IV - Other Interesting Features in the Data.ipynb.

5 Calculating Accuracies and Uncertainties

The predictions in sections 2 and 3 were binary. For example, the predictions in section 2 were whether a house was charging or not, which could be labeled as 0 for not charging and 1 for charging. The predictions in section 3 were whether a house owned an EV, which could be labeled as 0 for no EV ownership and 1 for EV ownership. If the classifier outputted 0 and the true result was 0, this would be labeled as a True Negative (or TN). There exist definitions for all four possibilities, as described in the table below.

Prediction	True Value	Name	Abbreviation
0	0	True Negative	TN
0	1	False Negative	FN
1	0	False Positive	FP
1	1	True Positive	TP

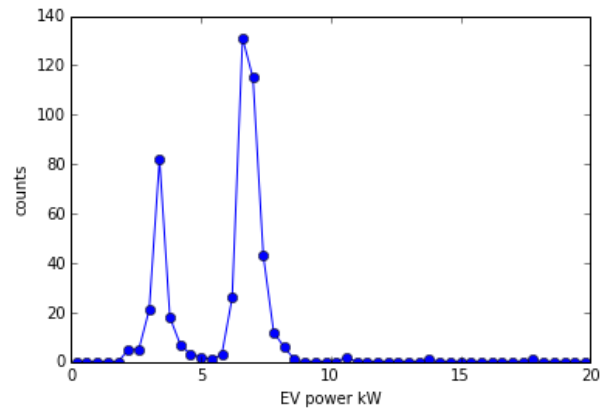


Figure 3: Histogram of EV power consumption by household.

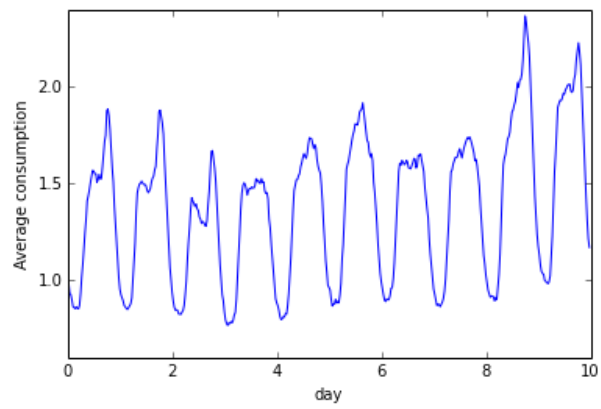


Figure 4: Average household consumption over time.

The accuracy can be calculated using

$$\text{Accuracy} = (\text{TN} + \text{TP})/(\text{TN} + \text{TP} + \text{FN} + \text{FP}). \quad (1)$$

Other types of accuracies, such as the Positive Accuracy Rate and Negative Accuracy Rate, are calculated in Part II and Part III of the notebooks.

Since each of these four quantities is a count, the uncertainties were assumed to be Poissonian, which were calculated by taking the square root of each count. The uncertainties were then propagated using partial derivatives.